

ClaimSure



Claims Fraud Triage Copilot

A Conversational AI agent for claims fraud detection, developed through Snowflake Cortex.

Snowflake Cortex

AI engine for intelligent code execution & LLM reasoning

Streamlit in Snowflake

Interactive conversational UI built directly inside Snowflake

Claims Fraud Detection

Automated triage, anomaly scoring & policy inspection

TEAM NAME

SnowNexus

TEAM MEMBERS

Aniket Jain • Akshaya Jayakumar

Udit Saini • Peruri Mohana Satya

 Claims Fraud Triage Copilot | Powered by Snowflake Cortex

Problem Statement & Why It Matters



THE CHALLENGE TODAY



Data Locked in Tables

Claim amount, fraud score, and policy status sit in database tables business users can't query directly.



Contracts Nobody Rereads

Coverage and exclusion clauses live in long PDFs nobody has time to fully re-read per claim.



Unknown Data Trust

Is the underlying data reliable this week? Usually only the data engineering team knows.



No Risk Triage

Every claim is queued equally, regardless of risk — no way to fast-track the low-risk majority.

Why

It

Matters

By unifying claims data, policy documents, and data-quality signals into a single Cortex Agent, adjusters and business users can investigate fraud using plain English. This eliminates the need for SQL queries, manual PDF reviews, or waiting on data engineering to verify data accuracy. The result is faster payouts for clean claims, reduced risk of missed exclusions, and a scalable Cortex Agents pattern that can be expanded to other claims workflows without additional infrastructure.

 *Every minute spent re-reading a policy is a minute not spent on the claim that actually needs it.*

Proposed Solution - AI Powered Copilot



A unified AI assistant built on Snowflake Cortex features that helps insurance claims adjusters fast-track honest claims and focus investigation on suspicious ones.

01 · Cortex Analyst

Self-Service Analytics

Ask business questions in plain English — no SQL or BI dashboard needed.

02 · Cortex Search

Document Intelligence

Ask what a policy covers or excludes — answered from the actual contract text.

03 · Cortex Complete

Data Quality Root Cause

Ask why a metric looks off — get a plain-English explanation, not just a score.

How It Works




 **One agent, three Cortex specialists** — not three separate tools. An orchestrating Cortex Agent reads each question and routes it automatically.

WHO IS THIS FOR?




Claimsure, One platform. Six teams.


Each reading the same data through the lens of their own job —
and no new tool for any of them to learn.



Architecture – End-to-End Data Flow



Entity Relationship Diagram — Insurance AI Hub



PK Primary Key | FK Foreign Key | Solid lines denote direct parent-child relationships; dashed lines indicate cross-domain references.

Old Approach vs New Approach



✗ OLD WAY — Without ClaimSure

- ✗ Manually re-reading policy PDFs per claim
- ✗ Business users can't query claims data directly
- ✗ All claims queued equally, regardless of risk
- ✗ Data quality checked only by data engineering
- ✗ High call-center volume on “what's covered”
- ✗ A single static fraud score, no context

✓ NEW WAY — With ClaimSure

- ✓ Instant exclusion-clause lookup via Cortex Search
- ✓ Plain-English analytics via Cortex Analyst
- ✓ Low-risk claims fast-tracked automatically
- ✓ Conversational root-cause explanations for anyone
- ✓ Self-serve coverage answers, fewer calls
- ✓ Reasoned verdict combining score, history & policy text

⚡ **ClaimSure doesn't just speed up lookups** — it changes how claims get triaged entirely.

Expected Outcomes & Measurable Impact

Quantifiable business value delivered by AI-driven claims triage



80%

Claims Fast-Track

Low-risk claims identified instantly by AI — skip full manual review

100%

Consistency

Every claim gets the same checks every time, regardless of workload

0

New Infrastructure

Built entirely on Snowflake Cortex, Streamlit, and existing warehouse data

~90%

Triage Time Saved

Seconds for AI verdict vs hours of manual cross-referencing across systems

◆ Fraud Detection Accuracy

AI-scored fraud probability enables data-driven investigation prioritization.

◆ Cost Savings

Reduce fraud investigation team scaling needs — AI handles triage at any volume.

◆ Customer Satisfaction

Faster payouts for genuine claims — 80% processed without full manual review.

◆ Data Trust

Built-in Data Quality monitoring ensures AI verdicts are based on reliable, timely data.

Streamlit APP User Interface (Agent)

The interface: one chat, every question type

Question asked: "How's our claims data quality trending?"





CLAIMS FRAUD TRIAGE COPILOT

One AI assistant that reads claims, policies, and documents together.

Powered by

- Cortex Agent (orchestration)
- Cortex Analyst (structured data)
- Cortex Search (policy docs)

ClaimSure Agent Claims Dashboard Data Quality Dashboard

ClaimSure Agent

Ask about any claim, policy, or customer. The Cortex Agent queries structured data and policy documents simultaneously, then synthesizes a risk verdict with cited evidence.

Try a quick question:

Is claim CLM-00050 suspicious?Top 10 highest fraud score open claimsWhat exclusions apply to auto policies?Claims breakdown by type and avg fraud score

Type your question here and press Enter or click Ask...

Ask

CONVERSATION

YOU

How's our claims data quality trending?

CLAIMSURE COPILOT


I'll look into your claims data to assess quality trends. Here's the read on claims data quality over the trailing 12 months (Sep 2025 – Sep 2026, based on `claim\_date`).

Structural integrity is clean and stable — no defects trending in any month:

- 0 orphan claims — every claim links to a valid policy
- 0 missing claim amounts, claim statuses, or fraud scores

Streamlit APP User Interface (Dashboard)

Claims Dashboard



Streamlit APP User Interface (Dashboard)

Data Quality Dashboard



One AI assistant that reads claims, policies, and documents together.

Powered by

- Cortex Agent (orchestration)
- Cortex Analyst (structured data)
- Cortex Search (policy docs)

ClaimSure Agent Claims Dashboard Data Quality Dashboard

Data Quality Dashboard

Trust layer for the AI copilot — tracks completeness, accuracy, consistency, and timeliness across all source tables so verdicts are grounded in reliable data.

Last refreshed: Sep 23, 2026 01:25:56 AM

Latest Scores by Table



DQ Score Trend Over Time



Streamlit APP User Interface

Main KPI Metrics



AT A GLANCE

Main metrics from the live dashboards

Claims Dashboard



Data Quality Dashboard — latest scores by table



Overall data quality holds a healthy 88.5% average across all seven tables, with CUSTOMERS the one area flagged for improvement.

Worked Example — From Question to Verdict



“Claim CLM-4521 — is this a valid claim, and is anything about it suspicious?”

What Cortex Analyst Pulls Back

Claim type - Water damage — residential

Claim amount - **\$18,400**

Fraud score - **0.71 (elevated)**

Policy status - Active, valid coverage window

Claims, last 6 months - 2 prior claims (same customer)

What Cortex Search Finds

Retrieved clause — Section 4.2, Exclusions:

“Water damage originating from gradual seepage is not covered under this policy.”

Matched from POLICY_DOCUMENTS via Cortex Search over DOCUMENT_CHUNKS
— no manual PDF reading required.

HIGH RISK

- Claimed damage (gradual seepage) falls under a documented policy exclusion
- Customer has filed 2 similar claims in the past 6 months
- Fraud score of 0.71 is above the elevated-risk threshold

Adjuster action: escalate to fraud investigation queue — skip fast-track approval.

 **Plain-English risk verdict**, with the evidence attached — in seconds.

Discussion & Next Steps



Any Questions?

Questions, feedback, and next-step alignment welcome.



Questions

Deep-dive into Snowflake Cortex architecture, agent routing, or data quality controls.



Feedback

Share insights on triage accuracy, fraud scoring thresholds, and adjuster workflow ergonomics.



Next Steps

Align on pilot scope, timeline, MCP system integrations, and row-level security setup.

